AI Customer Complaint Triage with n8n and Google Gemini

This project is an AI-powered customer complaint triage workflow built in n8n, a visual workflow automation platform, using Google Gemini as the reasoning engine. It shows how a large language model can be placed inside a real business process, so that it makes decisions and moves data between systems rather than only chatting.

The workflow starts with a Webhook node that accepts a POST request containing the customer's name and complaint. In the test run, sent from Postman, the customer reported that a UPI payment had failed but 2000 rupees had still been debited. The complaint is passed through an n8n expression into a Basic LLM Chain, which is connected to the Gemini 3.1 Flash Lite chat model. The prompt tells the model to return only four fields: Category, Priority, Reason and a Customer Response. Gemini classified the ticket as a Payment/Transaction Issue with High priority and wrote a polite reply asking for the transaction ID, using only 109 tokens.

An If node then checks whether the model output contains "Priority: High". On the true branch, an Edit Fields node builds a structured ticket with the status ESCALATED, priority HIGH and an action routing it to the payment support team. Finally, a Respond to Webhook node returns a JSON response in the same HTTP request, so the workflow behaves like a live API endpoint.

The project demonstrates LLM-based classification, conditional routing, data mapping and synchronous webhook responses, with credentials kept securely in n8n. Planned improvements include structured JSON output, a Switch node for medium and low priority, dynamic responses, ticket logging and webhook authentication.